

Backend Technical Interview

Context

You are a Software Engineer at a company that sells datasets to Quantitative Traders so that they can backtest their strategies on historical pricing data. You are responsible for building the Backend for the **Entitlements** part of the system (i.e., providing customers with programmatic access to datasets).

The 2 main users of this part of the system are internal Operations team members (**Ops**) and external Quantitative Trader customers (**Quant**). The typical workflow consists of a Quant browsing the datasets available on the platform, and requesting access to one or more datasets; then based on the validity of the request, it is either approved/rejected automatically, or it notifies an Ops who triggers the approval or refusal of the request; finally a Quant is notified of the decision, and - if approved - is able to fetch the content of the dataset for a given frequency.

To get the product in the hands of users and get feedback early, you decide to build an MVP. This means your implementation should work end-to-end, but be pragmatic in terms of the functionalities implemented, while still considering what would be needed for a production-ready system at scale.

Scope

Given the time constraints and MVP nature of the exercise, the following functionalities will have a **limited scope** or be **out of scope**:

- **Datasets**: Each dataset will be for 1 crypto asset, and will have a **name** and a **symbol**, and be available in various **frequencies** (see table below). The datasets info will be publically available, and the pricing data will come from this [CoinCap API endpoint](#).

Only seed the datasets below

Name	Symbol	Frequencies
Bitcoin	BTC	1 hour + 1 day + 1 month
Ethereum	ETH	1 day

- **Users**: Users should only be represented by an API Key with a role attached to it (either seeded or created via API).
- **Payment**: The dataset access request approvals/refusals are at the discretion of the Ops and not dependent on a Quant buying a package/subscription.
- **Notification**: Some functionalities will be asynchronous by nature (e.g., decision notification) and can be implemented using the pattern that you wish (e.g., background job, polling, websocket).

Exercise 1: System Design

Using a digital white-boarding tool of your choice:

1. Draw out the technical components of the system and their interactions
2. Highlight the underlying cloud services you would use to deploy it to a production environment
3. Break out problem-space into subdomain (in a Domain Driven Design fashion), which can be implemented as microservices
4. Define the database modelling

Exercise 2: Code Implementation

Think of the following requirements as **Behavioural Driven Development test-cases** - they all need to pass for the product to be considered working. The details of implementation are up to you and will be discussed during your presentation.

Functional

Mandatory

- ☐ A Quant can view the metadata (i.e., name & symbol) and available frequencies for on all datasets on the platform (i.e., BTC & ETH)
- ☐ A Quant can request access to 1 or more frequencies for a given dataset
- ☐ An Ops can view all pending dataset access requests
- ☐ An Ops can approve or refuse a request, give access (or not) to a dataset for a frequency to the Quant
- ☐ A Quant can view the pricing data (in USD) of the datasets/frequencies they have access to

Optional

- ☐ Backfill the market-cap (in USD) of each dataset every day at 8pm UTC (using this [CoinCap API endpoint](#))
- ☐ Implement a way for Ops to provide access to datasets/frequencies for a limited amount of time (e.g., trial) and revoke accesses

Non-functional

Mandatory

- ☐ Use NodeJS with Typescript
- ☐ Use a Postgres database
- ☐ Use REST for the API layer
- ☐ Dockerise the application and run it locally using docker-compose
- ☐ Write tests to cover the critical path

Optional

- ☐ Use the NestJS framework
- ☐ Use the Sequelize ORM
- ☐ Implement at least 2 micro-services